


Prompt Engineering for Algorithmic Code Generation

Comparative Analysis of Reasoning Strategies on LLM Performance

Team Members:

- Prakriti Adhikari
- Fiza Ashraf
- Aymen Noor



Context & Problem Definition


Problem Statement (Text)

Given an array, find the maximum subarray sum. The code prenotes givere the array, find the maximum function find the maximum subarray sum. If he constraints t: time to the maxurimum subarray Sum was to anything.

Constraints:

$N \leq 10^5$
 $N \leq 10^5$
...

Large Language Model (LLM)



Refinement & Debugging

Generated Source Code

```
def max_subarray_sum(arr):...  
    for i in instanster:  
        so_far = arr[0]  
        if irr; i in lenght.array_rum0: arr+):...  
            ll := max_subarray_sum()  
            return max_so_far
```

Evaluation & Testing
(Sample Cases)

Fails/Errors

Context & Problem Definition

The Challenge:

- Large Language Models (LLMs) often **fail in highly reasoning-intensive domains** like competitive programming.
- These problems require **complex, multi-step logical deduction** and robust algorithm selection.

Project Task:

- Investigate and systematically compare the effect of **various prompt engineering strategies** on LLM performance when solving complex algorithmic problems.

Real-World Relevance:


- Improving LLM capability for **advanced developer tooling**, automated code review, and high-stakes technical domains

Methodology - Pipeline & Dataset

Dataset:

- **APPS (Automated Programming Problem Set)** dataset.
- Subset: 100 problems, all tagged with "**Interview**" difficulty.
- Evaluation is **objective**, based on Ground Truth Test Cases.

System Design Flowchart: EXPERIMENTAL PIPELINE: HIGH-THROUGHPUT LLM EVALUATION ON APPS DATASET



Experimental Pipeline (6 Stages):

1. **Data Ingestion:** Load problem descriptions and ground truth test cases.
2. **Prompt Generation:** Insert problem into one of **seven prompt templates**.
3. **LLM Inference:** Send prompt to primary model (**Gemini 2.5 Flash Lite API**).
4. **Output Parsing:** Extract reasoning ([BEGIN REASONING] tags) and code (python block).
5. **Secure Execution & Evaluation:** Execute code in a subprocess against test cases with a **strict 5s timeout**.
6. **Logging & Completion:** Log result as **PASS, FAIL, or ERROR**.

Prompting Strategies Compared


Strategy	Primary Goal	Key Feature
Zero-Shot	Assess inherent ability with mandated reasoning structure.	No examples. Mandates detailed reasoning steps.
Zero-Shot CoT	Zero-Shot variant with slightly different mandated steps.	Mandates steps: Restate problem, Identify constraints, Derive optimal solution.
Few-Shot	Evaluate benefit of in-context learning, isolating from reasoning.	Prepends 3 examples (problem/solution), no reasoning chains .
Few-Shot CoT	Assess maximum benefit of Gold Standard CoT .	Prepends 3 examples with full, high-quality manual reasoning and solution.
Auto-Few-Shot CoT	Evaluate scalability of self-generated reasoning .	LLM <i>first</i> generates reasoning for the 3 examples, which are then used.

Key Finding 1: Model Head-to-Head

Gemini 2.5 Flash Lite Consistently Outperformed Llama 4

Metric	Gemini 2.5 Flash Lite (Total)	Llama 4 (Total)	Interpretation
Total PASS Results	325	137	Gemini is generally more capable of generating logically correct solutions.
Total FAIL Results (Logic Mismatch)	268	361	Llama struggled significantly more with logical correctness .
Total ERROR Results	100	102	Both models exhibited similar rates of runtime instability.

Model Performance Comparison: Gemini 2.5 Flash Lite vs. Llama 4




Key Finding 2: Strategy Effectiveness

Gemini: Simplicity Outperforms Complexity

Strategy	Gemini Accuracy (%)	Llama Accuracy (%)
Zero-Shot	55.77	35.00
Zero-Shot CoT	54.81	28.00
Analytic CoT	52.88	N/A
Few-Shot CoT	38.46	15.00
Auto-Few-Shot CoT	33.33	19.00

Gemini vs. Llama: Model Accuracy per Strategy



Slide 6: Key Finding 2: Strategy Effectiveness

Interpretation for Gemini: Providing direct instructions and a structured analytical path (Zero-Shot CoT) was most effective.

Complex Few-Shot strategies (especially Auto-CoT) actively degraded performance due to context size issues and parsing variability.

Llama 4: CoT variants did not provide a clear, consistent accuracy advantage, and the model struggled severely with Few-Shot variants

Failure Analysis

Logic Mismatch (FAIL):

- **Overwhelming primary cause** of non-PASS outcomes for both models (Llama: 361; Gemini: 268).
- Confirms the challenge is in **algorithmic derivation and implementation accuracy**, not syntax.

Instability and Efficiency:

Timeout Errors: Suggests efficiency issues—code is correct but **too slow** for the 5-second limit (Llama: 36; Gemini: 21).

Auto-Few-Shot CoT was particularly unstable for Gemini, showing a **39% ERROR rate**, highlighting instability from model-generated training context.

- **Qualitative Insight:** Successful Zero-Shot CoT runs on Gemini often produced **concise, correct algorithmic reasoning** (e.g., using dynamic programming/prefix array), indicating the prompt successfully compels the articulation of the optimal approach.

Conclusion

Simpler prompting strategies consistently outperformed complex Few-Shot CoT variants for the high-performing model (Gemini).

Relying on the model's inherent, structured reasoning ability is more robust than introducing excessive, potentially error-prone context.

Limitations:

- Persistent dominance of **Mismatch (logic errors)**.
- **Strict 5-second timeout** led to performance issues from unoptimized solutions.
- Limited to 100 "Interview" difficulty problems.

Future Work:

- **Enhancing Logical Correctness:** Focus on **Self-Correction (Self-Refine)** mechanisms, where the model iteratively debugs its code using unit test feedback.
- **Model Fine-Tuning:** Fine-tuning a base model (e.g., Llama 4) on the APPS dataset to drastically improve logical correctness and output reliability.